ResiDual for Audio: Spectral Reweighting of Residual Streams in CLAP Models

April 7, 2026

Lorenzo Arcioni

Abstract

This is a LaTeXtemplate for writing your project report, to be submitted as part of the final exam. The template can not be modified (you can not change margins, spaces, etc.), and using this template is mandatory. Please read the main text for further details.

A. CLAP and HTS-AT: Architecture, Notation, and Residual Decomposition

This appendix is a self-contained reference for the two components at the core of our system. We begin by describing the CLAP dual-encoder model and its configuration (§A.1). We then present the HTS-AT audio backbone in detail, covering input preprocessing, hierarchical stage structure, and the internal mechanics of each attention block (§A.2). Finally, we derive the per-head residual decomposition that underlies our analysis and define the aggregated representations used throughout the paper (§A.3). All notation is introduced inline and collected in Table 5 for reference.

A.1. CLAP Overview

CLAP (Contrastive Language–Audio Pretraining) (**?**) is a dual-encoder model that aligns audio and text representations in a shared embedding space via contrastive learning. The audio encoder is HTS-AT (**?**) and the text encoder is GPT-2 (**?**) (embedding dimension 768), whose weights are frozen throughout training. Both encoders are independently projected to a joint space of dimension $d_{\text{proj}} = 1024$ via dedicated linear projection heads; similarity is measured by temperature-scaled cosine similarity with InfoNCE loss at temperature $\tau = 0.003$. Zero-shot audio classification is performed by comparing an audio embedding against the embeddings of textual class prompts.

The full configuration used in this work is reported in Table 3.

Table 1. CLAP configuration parameters.

Parameter	**Value**
Text encoder	GPT-2
Text encoder embedding dim	768
Audio encoder	HTS-AT
Audio encoder output dim (D_3)	768
Joint projection dim (d_{proj})	1024
Contrastive temperature τ	0.003
Sampling rate	44 100 Hz
Audio duration	7 s
Mel bands	64
FFT window	1024
Hop size	320
$f_{\min}$ / $f_{\max}$	50 Hz / 8 000 Hz

A.2. HTS-AT Architecture

HTS-AT (**?**) is a Swin-Transformer variant adapted for audio spectrograms. Its computation is organised into four hierarchical *stages*, each composed of several *blocks*, as illustrated in Figure 6. We describe the pipeline from raw audio to the final embedding in order: input preprocessing, patch embedding, stage structure, and the attention mechanism inside each block.

Input preprocessing and patch embedding. Each audio clip is converted to a 64-band log-mel spectrogram ($f_{\min} = 50$ Hz, $f_{\max} = 8{,}000$ Hz, STFT window 1024, hop 320) and normalised per mel-band. The resulting time–frequency matrix is rearranged into a square image $\mathbf{x}_{\text{img}} \in \mathbb{R}^{1\times 256\times 256}$ by folding the time axis into four contiguous segments of 256 frames stacked vertically over the 64 mel-frequency bands ($4\times 64 = 256$ rows, 256 columns). This image is then split into 4×4 non-overlapping patches via a strided `Conv2d` layer (kernel 4×4, stride 4), followed

Email: Lorenzo Arcioni <arcioni.1885377@studenti.uniroma1.it>.

Deep Learning and Applied AI 2025, Sapienza University of Rome, 2nd semester a.y. 2024/2025.

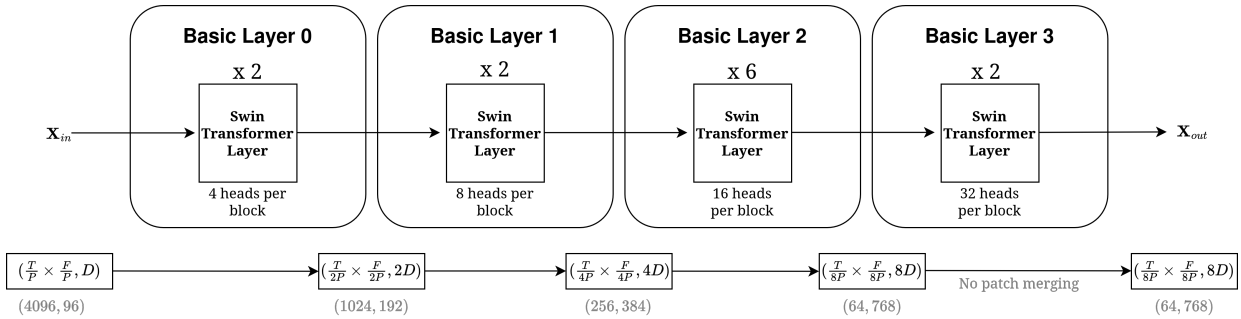

Figure 1. HTS-AT hierarchical architecture. The model consists of four stages of increasing embedding dimension: Stage 0 (2 blocks, 4 heads), Stage 1 (2 blocks, 8 heads), Stage 2 (6 blocks, 16 heads), and Stage 3 (2 blocks, 32 heads). PatchMerging is applied after Stages 0, 1, and 2, halving the spatial side length and doubling the feature dimension at each transition. Stage 3 has no PatchMerging. The input token sequence has spatial size $64 \times 64 = 4{,}096$ tokens with embedding dimension $D_0 = 96$; after three PatchMerging operations the final spatial size is $8 \times 8 = 64$ tokens with $D_3 = 768$.

by `LayerNorm`, yielding the input token sequence

$$\mathbf{Z}^{(0,0)} = \mathbf{X}_{\text{in}} \in \mathbb{R}^{4096 \times 96}, \tag{1}$$

where $4096 = 64 \times 64$ tokens each have embedding dimension $D_0 = 96$.

Stage structure and PatchMerging. The four stages operate at progressively coarser spatial resolutions, as summarised in Table 4. We index the residual stream as $\mathbf{Z}^{(\ell,b)}$, where $\ell \in \{0, 1, 2, 3\}$ is the stage index and $b \in \{1, \ldots, B_\ell\}$ is the block index within that stage.

After Stages 0, 1, and 2, a *PatchMerging* layer concatenates each 2×2 neighbourhood of spatially adjacent tokens into a single vector of dimension $4D_\ell$, then projects it to $2D_\ell$ via a bias-free linear layer. This halves the spatial side length S_ℓ and doubles the embedding dimension:

$$\mathbb{R}^{S_\ell^2 \times D_\ell} \xrightarrow{\text{PatchMerging}} \mathbb{R}^{(S_\ell/2)^2 \times 2D_\ell}. \tag{2}$$

Stage 3 has no PatchMerging. The resulting spatial side lengths are $S_\ell \in \{64, 32, 16, 8\}$ and the corresponding embedding dimensions are $D_\ell \in \{96, 192, 384, 768\}$ (see Table 4).

Block computation: W-MSA and MLP. Each block b at stage ℓ applies two sub-layers in sequence, both preceded by LayerNorm and connected via residual additions:

$$\mathbf{Z}^{(\ell,b)} \leftarrow \mathbf{Z}^{(\ell,b-1)} + \text{W-MSA}_{\ell,b}\Big(\text{LN}\Big(\mathbf{Z}^{(\ell,b-1)}\Big)\Big), \tag{3}$$

$$\mathbf{Z}^{(\ell,b)} \leftarrow \mathbf{Z}^{(\ell,b)} + \text{MLP}_{\ell,b}\Big(\text{LN}\Big(\mathbf{Z}^{(\ell,b)}\Big)\Big). \tag{4}$$

Even-indexed blocks use standard Window Multi-head Self-Attention (W-MSA); odd-indexed blocks use Shifted-Window MSA (SW-MSA), which shifts the partition by $(w/2, w/2)$ tokens to enable cross-window interactions. At Stage 3, where $S_3 = w = 8$ so the grid coincides with a single window, the shift degenerates to zero and all blocks use W-MSA. The MLP sub-layer is a two-layer feed-forward network with hidden dimension $4D_\ell$ and GELU activation.

WindowAttention in detail. The W-MSA module at block (ℓ, b) first partitions the S_ℓ^2 tokens into $N_w^\ell = S_\ell^2/M$ non-overlapping windows of $M = 64$ tokens each, then applies multi-head attention independently within each window. Queries, keys, and values are computed via a single fused projection $W_{\ell,b}^{QKV} \in \mathbb{R}^{D_\ell \times 3D_\ell}$:

$$\text{LN}(\mathbf{Z}^{(\ell,b-1)})\, W_{\ell,b}^{QKV} \xrightarrow{\text{split}(D_\ell)} \mathbf{Q},\, \mathbf{K},\, \mathbf{V} \in \mathbb{R}^{N_w^\ell \times M \times D_\ell}. \tag{5}$$

Each tensor is then reshaped to isolate the H_ℓ heads, yielding per-head slices $\mathbf{Q}_h, \mathbf{K}_h, \mathbf{V}_h \in \mathbb{R}^{N_w^\ell \times M \times d_h}$ with $d_h = D_\ell/H_\ell = 24$. In the SW-MSA variant, a cyclic shift is applied before window partitioning and reversed afterward.

For each head $h \in \{1, \ldots, H_\ell\}$, attention is computed as:

$$\mathbf{H}_{\ell,b,h} = \text{Softmax}\left(\frac{\mathbf{Q}_h \mathbf{K}_h^\top}{\sqrt{d_h}} + \mathbf{B}_{\ell,b,h}\right) \mathbf{V}_h \in \mathbb{R}^{N_w^\ell \times M \times d_h}, \tag{6}$$

where $\mathbf{B}_{\ell,b,h} \in \mathbb{R}^{M \times M}$ is the learned relative position bias for head h of block (ℓ, b). This bias encodes the relative spatial offset between each pair of tokens *within* a window; it is shared across all N_w^ℓ windows (broadcast) and is independent for each head. Concretely, it is read from a parameter table of shape $((2w-1)^2, H_\ell) = (225, H_\ell)$ via a fixed index mapping, so each head h has its own column of learnable values.

All H_ℓ head outputs are then concatenated and passed through the block-specific output projection $W_{\ell,b}^O \in$

Table 2. HTS-AT stage-level architectural parameters. The per-head dimension $d_h = D_\ell/H_\ell = 24$ is constant across all stages. S_ℓ is the spatial side length of the token grid at stage ℓ; $N_w^\ell = S_\ell^2/M$ is the number of attention windows, with $M = w^2 = 64$ tokens per window ($w = 8$).

Stage ℓ	**Blocks** B_ℓ	**Heads** H_ℓ	**Dim** D_ℓ	**Spatial** S_ℓ^2	**Windows** N_w^ℓ
0	2	4	96	64×64	64
1	2	8	192	32×32	16
2	6	16	384	16×16	4
3	2	32	768	8×8	1
Total heads H_{tot}	$2{\cdot}4 + 2{\cdot}8 + 6{\cdot}16 + 2{\cdot}32 = 8 + 16 + 96 + 64 = 184$				

$\mathbb{R}^{D_\ell \times D_\ell}$ with bias $\mathbf{b}_{\ell,b}^O \in \mathbb{R}^{D_\ell}$, giving the W-MSA output

$$\mathbf{A}_{\ell,b} = \left[\mathbf{H}_{\ell,b,1} \| \cdots \| \mathbf{H}_{\ell,b,H_\ell}\right] W_{\ell,b}^O + \mathbf{b}_{\ell,b}^O \;\in\; \mathbb{R}^{N_w^\ell \cdot M \times D_\ell}, \tag{7}$$

which is added to the residual stream in Eq. (15). For a complete step-by-step account of all tensor shapes across every stage, see https://github.com/lorenzo-arcioni/ResiDual-CLAP/blob/main/README.md.

A.3. Per-Head Residual Decomposition

The additive residual structure of Eq. (15) allows us to decompose the W-MSA contribution $\mathbf{A}_{\ell,b}$ exactly into independent per-head terms. Denoting by $W_{\ell,b,h}^O \in \mathbb{R}^{d_h \times D_\ell}$ the row slice of $W_{\ell,b}^O$ corresponding to head h (rows $[(h-1)d_h,\ hd_h)$), we distribute the output projection over heads:

$$\begin{aligned}\mathbf{A}_{\ell,b} &= \left[\mathbf{H}_{\ell,b,1} \| \cdots \| \mathbf{H}_{\ell,b,H_\ell}\right] W_{\ell,b}^O + \mathbf{b}_{\ell,b}^O \\ &= \sum_{h=1}^{H_\ell} \left(\mathbf{H}_{\ell,b,h}\, W_{\ell,b,h}^O + \tfrac{\mathbf{b}_{\ell,b}^O}{H_\ell} \right) \;=\; \sum_{h=1}^{H_\ell} \widehat{\mathbf{H}}_{\ell,b,h},\end{aligned} \tag{8}$$

where we define the *per-head projected contribution*

$$\widehat{\mathbf{H}}_{\ell,b,h} \;=\; \mathbf{H}_{\ell,b,h}\, W_{\ell,b,h}^O + \frac{\mathbf{b}_{\ell,b}^O}{H_\ell} \;\in\; \mathbb{R}^{N_w^\ell \times M \times D_\ell}. \tag{9}$$

This decomposition is exact and follows from the linearity of matrix multiplication: because $W_{\ell,b}^O$ acts on the concatenation of all head outputs, each head h effectively multiplies only its own row slice $W_{\ell,b,h}^O$, and the output bias can be distributed equally among heads without any approximation.

The decomposition holds within a single stage, where D_ℓ is constant. Across stage boundaries, PatchMerging changes both spatial resolution and embedding dimension, breaking any global additive structure; cross-stage comparisons must therefore account for this.

B. Post-Projection Head Analysis

The main analysis operates on *raw pre-projection* head outputs $\mathbf{H}_{\ell,b,h}$, which characterise the intrinsic computational geometry of each head independently of W^O. Here we describe a complementary analysis of the *projected* contributions $\widehat{\mathbf{H}}_{\ell,b,h}$, which reflect each head's additive influence on the residual stream after the output projection.

B.1. Motivation and Methodological Differences

The per-head projected contribution $\widehat{\mathbf{H}}_{\ell,b,h} \in \mathbb{R}^{N_w^\ell \times M \times D_\ell}$ is obtained by applying the row slice $W_{\ell,b,h}^O$ of the output projection to the raw head output, and distributing the output bias equally across heads (see Appendix A.3 for the exact decomposition). It lives in the full residual-stream space $\mathbb{R}^{D_\ell}$ of stage ℓ, offering a complementary analytical perspective:

1. **Residual-stream interpretability.** $\widehat{\mathbf{H}}_{\ell,b,h}$ is the quantity that each head writes additively to $\mathbf{Z}^{(\ell,b)}$ via Eq. (1). Its geometry therefore directly characterises each head's contribution to the representations that are consumed by downstream layers and, ultimately, the CLAP projection head P.

2. **Role of W^O.** Comparing pre- and post-projection dimensionality estimates reveals how the output projection transforms each head's internal computation: if dimensionality increases after projection, W^O expands the head's representational geometry in the residual stream; if it decreases, W^O compresses or rotates it.

B.2. Extraction Procedure

The same forward hooks described in Section 3.2 are reused. After intercepting $\mathbf{H}_{\ell,b,h}$ of shape (N_w^ℓ, M, d_h), we reconstruct the projected contribution by applying the corresponding row slice $W_{\ell,b,h}^O \in \mathbb{R}^{d_h \times D_\ell}$ retrieved from `self.proj.weight`, and adding the distributed bias

Table 3. Summary of notation used throughout the paper.

Symbol	Definition	Values / Notes
$\ell \in \{0, 1, 2, 3\}$	Stage index	
$b \in \{1, \ldots, B_\ell\}$	Block index within stage ℓ	$B_\ell \in \{2, 2, 6, 2\}$
$h \in \{1, \ldots, H_\ell\}$	Attention head index within block (ℓ, b)	$H_\ell = 4 \cdot 2^\ell \in \{4, 8, 16, 32\}$
S_ℓ	Spatial side length of the token grid at stage ℓ	$S_\ell \in \{64, 32, 16, 8\}$
$w = 8$	Attention window side length	
$M = w^2 = 64$	Tokens per attention window	
$N_w^\ell = S_\ell^2/M$	Number of attention windows at stage ℓ	$N_w^\ell \in \{64, 16, 4, 1\}$
$d_h = 24$	Per-head feature dimension (constant across stages)	$d_h = D_\ell/H_\ell$
$D_\ell = H_\ell \cdot d_h$	Total embedding dimension at stage ℓ	$D_\ell \in \{96, 192, 384, 768\}$
$d_{\text{proj}} = 1024$	CLAP joint embedding dimension	
$W_{\ell,b}^{QKV} \in \mathbb{R}^{D_\ell \times 3D_\ell}$	Fused QKV projection at block (ℓ, b)	
$W_{\ell,b}^{O} \in \mathbb{R}^{D_\ell \times D_\ell}$	Output projection at block (ℓ, b)	
$W_{\ell,b,h}^{O} \in \mathbb{R}^{d_h \times D_\ell}$	Row slice of $W_{\ell,b}^O$ for head h	rows $[(h-1)d_h,\ hd_h)$
$\mathbf{b}_{\ell,b}^{O} \in \mathbb{R}^{D_\ell}$	Output projection bias at block (ℓ, b)	
$\mathbf{B}_{\ell,b,h} \in \mathbb{R}^{M \times M}$	Relative position bias for head h at block (ℓ, b)	
$\mathbf{H}_{\ell,b,h} \in \mathbb{R}^{N_w^\ell \times M \times d_h}$	Raw head output at block (ℓ, b), head h (pre-$W_{\ell,b}^O$)	
$\widehat{\mathbf{H}}_{\ell,b,h} \in \mathbb{R}^{N_w^\ell \times M \times D_\ell}$	Per-head projected contribution (post-$W_{\ell,b}^O$); see Eq. (21)	
$\mathbf{r}_{\ell,b,h} \in \mathbb{R}^{d_h}$	Spatially aggregated raw head output; see Eq. (4)	
$\widehat{\mathbf{r}}_{\ell,b,h} \in \mathbb{R}^{D_\ell}$	Spatially aggregated projected head output; see Eq. (23)	
$\mathbf{R}_{\ell,b,h} \in \mathbb{R}^{n \times d_h}$	Dataset matrix of $\mathbf{r}_{\ell,b,h}$ across n samples	
$\widehat{\mathbf{R}}_{\ell,b,h} \in \mathbb{R}^{n \times D_\ell}$	Dataset matrix of $\widehat{\mathbf{r}}_{\ell,b,h}$ across n samples	
$P : \mathbb{R}^{768} \to \mathbb{R}^{1024}$	CLAP audio projection head (two-layer MLP with residual)	
n	Number of audio samples in the dataset	

Table 4. Comparison of the two analysis pipelines.

	Main (pre-W^O)	Appendix (post-W^O)
Representation	$\mathbf{H}_{\ell,b,h}$	$\widehat{\mathbf{H}}_{\ell,b,h}$
Ambient dim	$d_h = 24$ (constant)	$D_\ell \in \{96, 192, 384, 768\}$
Analysis space	$\mathbb{R}^{24}$ (native)	$\mathbb{R}^{D_\ell}$
Hook location	Before `.reshape` only	Before `.reshape`, then $\times W_h^O$
W^O applied	No	Yes (via W_h^O)
P applied	No	No
Captures	Internal attention computation	Contribution to residual stream

term $\mathbf{b}_{\ell,b}^O/H_\ell$, as in Eq. (20):

$$\widehat{\mathbf{H}}_{\ell,b,h} = \mathbf{H}_{\ell,b,h} W_{\ell,b,h}^O + \frac{\mathbf{b}_{\ell,b}^O}{H_\ell} \in \mathbb{R}^{N_w^\ell \times M \times D_\ell}. \quad (10)$$

We then spatially aggregate by mean-pooling over all windows and token positions:

$$\widehat{\mathbf{r}}_{\ell,b,h} = \frac{1}{N_w^\ell M} \sum_{i=1}^{N_w^\ell} \sum_{j=1}^{M} \widehat{\mathbf{H}}_{\ell,b,h}[i, j, :] \in \mathbb{R}^{D_\ell}, \quad (11)$$

yielding the dataset matrix $\widehat{\mathbf{R}}_{\ell,b,h} \in \mathbb{R}^{n \times D_\ell}$ by stacking across n samples.

B.3. Dimensionality Estimators in the Post-Projection Space

All estimators from Section 3.3 apply with $\widehat{\mathbf{R}}_{\ell,b,h}$ replacing $\mathbf{R}_{\ell,b,h}$. The key difference is that the ambient dimension now varies across stages: $D_\ell \in \{96, 192, 384, 768\}$.

PCA. The maximum number of non-trivial principal components of $\widehat{\mathbf{R}}_{\ell,b,h}$ is bounded by $\min(n, D_\ell)$, so a head at Stage 3 ($D_3 = 768$) has a strictly larger ambient ceiling than a head at Stage 0 ($D_0 = 96$). Direct comparison of explained-variance curves or intrinsic dimensionality estimates *across* stages is therefore confounded by this varying ceiling and must account for it explicitly; within a single stage, where D_ℓ is fixed, comparisons across blocks and heads are well-defined.

L/N Ratio. The ratio $d_{\text{PCA}_{99}}/d_{\text{TwoNN}}$ must be interpreted with awareness that the PCA ceiling grows with stage, which may inflate linear dimensionality estimates independently of any genuine increase in representational complexity.

B.4. Interpretation and Relationship to Main Results

Post-projection analysis is best suited to studying *how heads collectively shape the residual stream* and, ultimately, the CLAP embedding used for zero-shot classification. Pre-projection analysis—the perspective adopted in the main body—is best suited to studying *individual head specialisation* in isolation, free from the mixing introduced by W^O.

Concretely, if a head exhibits low intrinsic dimensionality in the pre-projection space but high dimensionality after projection, W^O acts as an expansion map for that head in the residual stream; the converse implies compression or selective routing of the head's computation into a lower-dimensional subspace of $\mathbb{R}^{D_\ell}$.

Limits of cross-stage analysis. A key difference between HTS-AT and isotropic vision transformers such as ViT (**?**) is that the ambient dimension of the residual stream is not constant: $D_\ell \in \{96, 192, 384, 768\}$ grows with each stage. As a consequence, the matrices $\widehat{\mathbf{R}}_{\ell,b,h}$ from different stages live in spaces of different dimension, and any dimensionality-reduction analysis—such as PCA—must be interpreted with care. Specifically, the maximum number of non-trivial principal components of $\widehat{\mathbf{R}}_{\ell,b,h}$ is bounded by $\min(n, D_\ell)$, so a head at Stage 3 ($D_3 = 768$) has a strictly larger ambient ceiling than a head at Stage 0 ($D_0 = 96$). Direct comparison of explained-variance curves or intrinsic dimensionality estimates across stages is therefore confounded by this varying ceiling, and must account for it explicitly. Within a single stage, where D_ℓ is fixed, comparisons across blocks and heads are well-defined and unambiguous.

C. Extended Dimensionality Analysis

This appendix provides comprehensive quantitative details and additional visualizations complementing the main results in Section 4.1.

C.1. Detailed Block-Level Statistics

Table 7 reports complete block-wise metrics across all 12 transformer blocks in HTS-AT, aggregating over the heads within each block as described in Section 3.3.

Relationship to Architecture. As illustrated in Figure 6, the spatial resolution decreases progressively through the network due to patch merging between stages. While this affects the number of tokens N processed by each attention head, our analysis focuses on the intrinsic dimensionality of the *head dimension* $d_h = 24$ after spatial aggregation (Eq. 4). Thus, the reported metrics characterize the semantic complexity of head representations independent of spatial resolution effects.

The hierarchical structure creates natural breakpoints for dimensionality analysis:

- **Stage 1 (Blocks 0–1)**: High spatial resolution ($T/2P \times F/2P$) but limited capacity ($D_0 = 96$). Early fusion of local spectral-temporal patterns.
- **Stage 2 (Blocks 2–3)**: First dimensionality jump coincides with 2× patch merging and head doubling. Transition from local to intermediate-scale features.
- **Stage 3 (Blocks 4–9)**: Deepest stage with 6 blocks enables iterative refinement at fixed spatial scale ($T/4P \times F/4P$) and capacity ($D_2 = 384$). Gradual dimensionality growth reflects progressive feature abstraction.
- **Stage 4 (Blocks 10–11)**: Maximum capacity ($D_3 = 768$) without further spatial reduction. Minimal dimensionality increase suggests saturation.

Table 5. Block-wise aggregated dimensionality metrics for TinySOL dataset. Blocks 0–1 (Stage 1), 2–3 (Stage 2), 4–9 (Stage 3), 10–11 (Stage 4). L = Linear ID ($d_{\mathrm{PCA}_{99}}$), N = Nonlinear ID (TwoNN), L/N = Linear-nonlinear ratio, EVR1 = First PC variance explained.

Block	Stage	Heads	L	N	L/N	EVR1
0	1	4	3.75	3.94	0.95	0.878
1	1	4	8.00	4.94	1.62	0.575
2	2	8	12.75	6.15	2.07	0.403
3	2	8	17.50	6.75	2.59	0.460
4	3	16	18.00	6.93	2.60	0.388
5	3	16	20.13	7.08	2.84	0.279
6	3	16	21.25	7.40	2.87	0.250
7	3	16	21.38	7.42	2.88	0.243
8	3	16	22.25	8.37	2.66	0.217
9	3	16	21.88	8.11	2.70	0.241
10	4	32	22.25	8.49	2.62	0.262
11	4	32	21.59	8.00	2.70	0.272

Interpretation.

- Block 0 operates near the linear regime (L/N $\approx$ 1), with almost 88% variance in the first PC, indicating extremely constrained early processing.
- The largest single-block jump occurs at the Stage 1→2 transition (blocks 1→2: ΔL = +4.75, +59%), corresponding to doubling of attention heads (4→8) and hidden dimension (96→192).
- Stage 3 exhibits gradual linear ID growth (18.00 → 22.25 over 6 blocks) despite constant architecture, suggesting intra-stage feature refinement through depth.

- Stage 4 shows minimal progression (blocks 10→11: ΔL = -0.66), consistent with representational saturation observed in the main text.

C.2. Extended Cross-Dataset Analysis

We replicate the full layer-wise analysis on ESC-50 and VocalSound to validate architectural generalizability. Tables 8 and 9 present complete statistics.

Table 6. Layer-wise dimensionality metrics for ESC-50 dataset (50 environmental sound classes, 1000 stratified samples).

Layer	$d_{\mathbf{PCA}_{99}}$	TwoNN	PR	EVR1
L0	5.2 ± 2.1	4.2 ± 0.7	1.8 ± 0.8	0.741 ± 0.187
L1	14.3 ± 2.8	6.2 ± 0.6	3.9 ± 1.3	0.449 ± 0.109
L2	20.1 ± 2.0	7.4 ± 0.8	7.5 ± 1.8	0.281 ± 0.081
L3	20.3 ± 1.2	7.8 ± 1.1	7.4 ± 1.9	0.279 ± 0.074

Table 7. Layer-wise dimensionality metrics for VocalSound dataset (6 vocal sound categories, 1000 stratified samples).

Layer	$d_{\mathbf{PCA}_{99}}$	TwoNN	PR	EVR1
L0	6.1 ± 2.3	4.6 ± 0.8	2.1 ± 1.0	0.698 ± 0.192
L1	16.2 ± 2.4	6.7 ± 0.5	4.5 ± 1.5	0.421 ± 0.117
L2	21.9 ± 1.9	8.0 ± 0.9	8.2 ± 2.0	0.263 ± 0.079
L3	22.7 ± 1.1	8.6 ± 0.8	8.3 ± 1.6	0.254 ± 0.071

Cross-Dataset Consistency Analysis. Despite differing semantic granularities (ESC-50: 50 classes, VocalSound: 6 classes, TinySOL: 14 classes), layer-wise trends remain remarkably stable:

- **L0 Concentration**: All datasets exhibit EVR1 ¿ 69% in Stage 1, confirming universal early spectral concentration.
- **L1 Expansion**: The L0→L1 dimensionality jump is consistent (TinySOL: +9.2, ESC-50: +9.1, VocalSound: +10.1 for $d_{\mathrm{PCA}_{99}}$), with coefficient of variation across datasets CV = 0.06.
- **L2-L3 Saturation**: All datasets show similar modest L2→L3 growth ($\Delta d < 2.0$), despite L3 having 2× the heads of L2, indicating architecture-driven capacity limits.
- **L/N Ratio Convergence**: By Stage 4, all datasets reach L/N ≈ 2.6–2.7, suggesting a universal nonlinear complexity regime independent of semantic domain.

C.3. Statistical Validation

C.3.1. ANOVA Results

One-way ANOVA tests for layer differences on TinySOL dataset:

Table 8. Statistical significance of layer effects on dimensionality metrics (TinySOL, $n = 184$ heads). All tests use $\alpha = 0.05$.

Metric	F-statistic	p-value	Significance
$d_{\mathrm{PCA}_{99}}$	262.64	< 0.001	***
TwoNN	58.93	< 0.001	***
PR	44.74	< 0.001	***
EffRank	76.03	< 0.001	***
EVR(PC1)	118.47	< 0.001	***

C.3.2. Post-Hoc Pairwise Comparisons

Bonferroni-corrected pairwise t-tests for $d_{\mathrm{PCA}_{99}}$ (6 comparisons, $\alpha_{\text{corrected}} = 0.0083$):

Table 9. Pairwise layer comparisons for linear intrinsic dimensionality (TinySOL). All comparisons significant at corrected $\alpha = 0.0083$.

Comparison	$\Delta d_{\mathbf{PCA}_{99}}$	Cohen's d	p-value
L0 vs L1	9.22	3.86	< 0.001
L0 vs L2	14.93	7.12	< 0.001
L0 vs L3	16.04	8.94	< 0.001
L1 vs L2	5.71	2.34	< 0.001
L1 vs L3	6.82	3.19	< 0.001
L2 vs L3	1.11	0.78	< 0.001

All effect sizes exceed Cohen's threshold for "large" effects ($d > 0.8$), with the L0 vs L3 comparison exhibiting extremely large effects ($d > 8$), confirming substantial representational differences across layers.

C.4. Additional Visualizations

C.4.1. PC1 Dominance and Boxplots

Figure 7 presents complementary views of dimensionality structure.

Observations from Panel (a):

- Only 2 heads in L0 (2.1%) fall below the 50% EVR1 threshold, compared to 89% of L3 heads, demonstrating near-universal early concentration.
- The EVR1 distribution shifts from unimodal (L0: concentrated near 0.7–0.8) to bimodal (L3: peaks at 0.2–0.3 and 0.35–0.4), suggesting emergence of head subpopulations with distinct specialization levels.

Observations from Panel (b):

- PR and EffRank exhibit parallel scaling, confirming their measurement of related spectral properties.
- TwoNN shows compressed scale relative to PCA99, visually emphasizing the linear-nonlinear gap discussed in the main text.

panel_E_pc1_dominance.png

(a) PC1 Variance Dominance

panel_F_boxplot.png

(b) Multi-Metric Boxplots

Figure 2. Extended dimensionality analysis. (a) First principal component variance explained across all 184 heads. Horizontal line at 50% marks equal-contribution threshold. Sharp decline from L0 (mean 73%) to L3 (mean 27%) quantifies transition from low-rank to distributed representations. (b) Boxplot comparison of four key metrics across layers, revealing consistent monotonic trends and increasing intra-layer variance in deeper stages (note wider boxes for L2-L3).

- Outliers (marked as individual points beyond whiskers) are rare in L0-L1 but frequent in L2-L3, consistent with increased head-level heterogeneity.

C.5. Computational Details

All analyses were performed on an NVIDIA A100 GPU (40GB) using PyTorch 2.0.1 and Python 3.10. Key implementation details:

- **Head Extraction**: Forward hooks registered via `torch.nn.Module.register_forward_hook`. Batch size 100 for extraction to balance memory and throughput.
- **PCA**: Computed via `sklearn.decomposition.PCA` with full SVD solver. Eigenvalue thresholding at machine epsilon ($\sim 10^{-7}$) to remove numerical noise.
- **TwoNN**: Implemented using `skdim.id.TwoNN` with default parameters (no k selection required).
- **MLE**: `skdim.id.MLE` with $k = 20$ neighbors, standard Euclidean metric.
- **Statistical Tests**: `scipy.stats` functions (`f_oneway`, `ttest_ind`, `spearmanr`) with standard settings.

Total extraction time: $\sim$45 minutes per dataset (1000 samples × 184 heads). Analysis pipeline code available at https://github.com/[ANONYMOUS]/residual-audio.

C.6. Reproducibility Checklist

To facilitate replication:

- Random seeds: 42 (Python), 42 (NumPy), 42 (PyTorch)
- CLAP version: `laion/clap-htsat-unfused` checkpoint from HuggingFace
- Audio preprocessing: CLAP default (64-band mel, 10s duration, 48kHz resampling)
- Dataset versions: ESC-50 v2.0, TinySOL v3.0, VocalSound official release
- Stratified sampling: `sklearn.model_selection.StratifiedShuffleSplit` with 1000 samples